

Off-policy value function approximation

Konpat Preechakul
Chulalongkorn University
September 2019

Two sided problems

- Off-policy target value

$$G_t \rightarrow p G_t$$

- Off-policy target state distribution

$$p^b \rightarrow p^\pi$$

Target value problem

Target value problem

- If the target is on-policy, it needs correction
 - MC, N-step need importance sampling
- Some targets are off-policy, no need for correction
 - Expected SARSA
 - Deterministic policies
 - Q-learning
 - Tree-backup

A bird eye view of on/off-policy

Algorithm	V/Q value	Make it off-policy	Variance	Bias
Monte Carlo	V	IS	High	Low
	Q	IS	High	Low
One-step SARSA	V	IS	Lower	High
	Q	IS	Lower	High
One-step Expected SARSA	V	IS	Lower	High
	Q	Already	Low	High
One-step TD with Deterministic Policy (including Q-learning)	V	IS	Lower	High
	Q	Already	Low	High
N-step SARSA (including lambda)	V	IS	Medium	Medium
	Q	IS	Medium	Medium
Tree backup	V	Already	Low	High
	Q	Already	Low	High

Semi-gradient with correction

- Off-policy N-step

$$g_{t:t+n} = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \gamma^n v(s_{t+n})$$

$$v(s_t) \leftarrow v(s_t) + \alpha \rho_{t:t+n-1} [G_{t:t+n} - v(s_t)]$$

- Off-policy N-step semi-gradient

$$\theta \leftarrow \theta + \alpha \rho_{t:t+n-1} (g_{t:t+n} - v_\theta(s_t)) \nabla_\theta v_\theta(s_t)$$

Q-learning with approximation

- No need for off-policy correction

$$q(s, a) \leftarrow$$

$$q(s, a) + \alpha [r(s, a) + \max_{a'} q(s', a') - q(s, a)]$$

for until q_θ is stable **do**

take action according to $q_\theta(s, a)$

collect (s, a, r, s')

$$\delta = r + \gamma \max_{a'} q_\theta(s', a') - q_\theta(s, a)$$

$$\theta \leftarrow \theta + \alpha \delta \nabla_\theta q_\theta(s, a)$$

end for

Summary


- Target value problem is straightforward
- Make sure that the target value is corrected
- Only **one part** of the problem

Target distribution problem

Target distribution problem

- Update distribution is “crucial” for convergence of semi-gradient
- Off-policy data = off-policy distribution
 - Convergence guarantee only on-policy distribution
- No convergence guarantee

Example of divergence



$$N = 10$$

update

$$N \leftarrow N + \alpha(2N - N) \cdot 1$$

$$10 + 0.1(10) \cdot 1$$

$$N \leftarrow 11$$


$$N \leftarrow N + \alpha(2N - N) \cdot 1$$

$$11 + 0.1(22 - 11) \cdot 1$$

$$N \leftarrow 12.1$$

frequent updates could diverge

On-policy is more reasonable



on-policy trajectory \longrightarrow


linear case : $2W \rightarrow 0$ in one-update ($\alpha=1$)

non-linear : not possible ! (needs multiple updates)


\hookrightarrow on-policy + non-linear
 \neq convergence

Intuitive divergence

- Problem of generalization of nearby states



Intuitive divergence



The deadly triad

- Off-policy
- Bootstrapping
- Approximation

Having the three at the same time causes
“instability”

(with semi-gradient)

Double is not deadly

Off-policy + bootstrap:

Q-learning

Off-policy + approximation:

Off-policy MC with approx.

Bootstrap + approximation:

On-policy linear TD

All are stable.

When does divergence happen?

- When you use off-policy data
- You don't correct the target distribution
 - Even you have corrected the target value
- You use approximation
- **You use semi-gradient**
 - Imply bootstrapping

Bandages for semi-gradient



The deadly triad bandages



- Off-policy => **more on-policy**
- Bootstrapping => **less bootstrap**
- Approximation => **less approximate**

More on-policy

- On-policy state distribution correction
- Importance sampling

$$\theta \leftarrow \theta + \alpha \rho_{0:t-1} \beta_t (r_{t+1} + \gamma v_{\theta}(s_{t+1}) - v_{\theta}(s_t)) \nabla_{\theta} v_{\theta}(s_t)$$

Distribution correction *Target correction*


$$\rho_{t:T-1} = \prod_{i=t}^{T-1} \frac{\pi(a_i | s_i)}{b(a_i | s_i)}$$

- High variance

Less bootstrapping

- Smaller discount
 - Smaller bootstrap
- N-step return
 - Smaller bootstrap
- Target networks
 - Realize (more) the independence assumption
- Loss constraints
 - Reducing dependency

Target network visualized



Target networks

$$\theta' \leftarrow \theta$$

for until q_θ is stable **do**

take action according to $q_\theta(s, a)$

collect (s, a, r, s')

$$\delta = r + \gamma \max_{a'} q_{\theta'}(s', a') - q_\theta(s, a)$$

$$\theta \leftarrow \theta + \alpha \delta \nabla_\theta q_\theta(s, a)$$

if every K steps **then**

$$\theta' \leftarrow \theta$$

end if

end for

how long do you need to get to $r + v'_{\theta'}$

Loss constraints

Reducing dependency between nearby states*

- Projected gradients

$$V_{\theta}(s) \uparrow \quad V_{\theta}(s') \uparrow$$

$$V_{\theta}(s) \uparrow \quad V_{\theta}(s') \text{ unchanged}$$

$$\text{we use } g \approx \nabla V_{\theta}(s) \text{ s.t. } g \perp \nabla V_{\theta}(s')$$


- Temporal consistency loss

$$TC(\theta) = \| V_{\theta}(s') - V_{\theta}(s') \|^2$$


Less approximate

- Increase the capacity of the approximator
 - This reduces the generalization effect
- First-order tabular update approximation
 - Intuition, table is stable even with off-policy

First-order tabular update approximation




slope problem



generalization problem

First-order tabular update approximation

- Problems from “generalization”
- Problems from “slope”
 - SGD => steepest
 - Tabular => proportional
- If each update “knows” its “generalization”, it could correct itself!



Why do we love semi-gradient so much?

- It is fast
- It gives favorable fixed-point if trained successfully
 - TD fixed point
- It is probably the one which is shown to work on large problems
 - Atari
 - Go

Moving to true gradient

Semi to true gradient

- **Semi-gradient is the culprit**
- True gradient guarantees convergence
 - Merit from SGD
 - Even in non-linear case
 - But to where?
- **Loss functions:**
 - TD error
 - Bellman error
 - Projected Bellman error

TD Error (TDE)

$$\text{TDE}(\theta) = \mathbb{E} \left[(R_{t+1} + \gamma v_{\theta}(S_{t+1}) - v_{\theta}(S_t))^2 | A_t \sim \pi \right]$$


$$\text{TDE}(\theta) = \mathbb{E} \left[\delta^2 | A_t \sim \pi \right]$$

True gradient:


$$\theta \leftarrow \theta + \alpha \delta_t [\nabla_{\theta} v_{\theta}(s_t) - \gamma \nabla_{\theta} v_{\theta}(s_{t+1})]$$

How good is the fixed point?

A-split problem




Semi-grad TD on A-split



$$\min_{v_\theta} \mathbb{E}_\pi [\delta^2]$$


TDE on A-split




$$\min_{\theta} \mathbb{E}_{\pi} [\delta^2] \quad \text{Tends to be large}$$

Temporal smoothing

Update visualized true SGD TD



Update visualized true SGD TD



**TDE seems bad,
Alternatives?**

Bellman error

- Bellman equation

$$v(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v(s')]$$

- Formulation for single state

$$\begin{aligned} \overline{\delta_\theta}(s) &= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\theta(s')] - v_\theta(s) \\ &= \mathbb{E}_\pi [\delta_s] \end{aligned}$$

Bellman error

$$\begin{aligned}\overline{\delta_\theta}(s) &= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\theta(s')] - v_\theta(s) \\ &= \mathbb{E}_\pi [\delta_s]\end{aligned}$$


- Weighted squared error for all states

$$\|\overline{\delta_\theta}\|_\pi^2 = \text{BE}(\theta) = \sum_s \mathbb{P}^\pi(s) [\mathbb{E}_\pi[\delta_s]^2]$$

↖ vs. $\mathbb{E}_\pi[\delta_s^2]$ TDE


- Minimize the above, using SGD

BE on A-split



$$\min_{\theta} \mathbb{E}_{\pi} [\delta]^2 \quad \text{Tends to be smaller}$$

BE on A-split (v2)



$$\min_{\theta} \mathbb{E}_{\pi} [\delta]^2 \quad \text{Tends to be smaller}$$

Doesn't help anymore

Indistinguishable state features

Bellman error's gradient

$$\overline{\delta_\theta}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\theta(s')] - v_\theta(s)$$

$$\|\overline{\delta_\theta}\|_\pi^2 = \text{BE}(\theta) = \sum_s \mathbb{P}^\pi(s) [\mathbb{E}_\pi[\delta_s]^2]$$

- Gradient

$$\nabla_\theta \text{BE}(\theta) = \sum_s \mathbb{P}^\pi(s) [\mathbb{E}_\pi[\delta_s] \mathbb{E}_\pi[\gamma \nabla_\theta v_\theta(s')] - \nabla_\theta v_\theta(s)]$$

Double sampling problem

$$\|\overline{\delta_\theta}\|_\pi^2 = \text{BE}(\theta) = \sum_s \mathbb{P}^\pi(s) [\mathbb{E}_\pi[\delta_s]^2]$$

$$\nabla_\theta \text{BE}(\theta) = \sum_s \mathbb{P}^\pi(s) [\mathbb{E}_\pi[\delta_s] \mathbb{E}_\pi[\gamma \nabla_\theta v_\theta(s') - \nabla_\theta v_\theta(s)]]$$

Bellman error summary

- Impractical
- Might converge to undesirable fixed point
- A better loss function needed
 - Projected Bellman error
 - How to get its gradients (GTD2)

Read further

- Sutton 2018, Chapter 11
- Maei, Hamid Reza. 2011. “Gradient Temporal-Difference Learning Algorithms.” University of Alberta.
- Topics:
 - Projected Bellman error
 - GTD2

Levels of guarantees

- **Stability vs convergence**
 - Stability is easier to guarantee
- **On-policy vs off-policy**
 - On-policy is easier to guarantee
- **Linear vs non-linear**
 - Linear is easier to guarantee

Convergence of off-policy non-linear function approximation:

- Doesn't imply a “good” fixed point

Summary of gradients

Semi gradients = converge only on linear with on-policy

- Semi-gradient TD
- You need luck and tricks

True gradients = converge even on non-linear, both on-policy and off-policy

- Value error
- TD error
- Bellman error

Assignment

- Q-learning with function approximation notebook (Github)